



ELL793

Computer Vision

Prof. Brejesh Lall

Assignment 3:

CIFAR-10

Dataset Classification

Submitted by:
Raghav Gupta (2017EE10544)
Prakhar Kanchan (2017EE30543)

Part 1

Goal

Train a CNN network with ResNet-18 as a backbone from scratch with CIFAR-10

Procedure

- Loaded the CIFAR-10 dataset, created a validation split and applied normalization on the same lines as the previous problem.
- The cell shown below describes the process of importing the dataset and pre-processing it for training.

```
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

dataset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=transform)
train_ds, val_ds = random_split(dataset, [40000, 10000])
test_ds = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=transform)
```

```
Train: X=(40000, 32, 32, 3), y=(40000, 10)
Val:   X=(10000, 32, 32, 3), y=(10000, 10)
Test:  X=(10000, 32, 32, 3), y=(10000, 10)
```

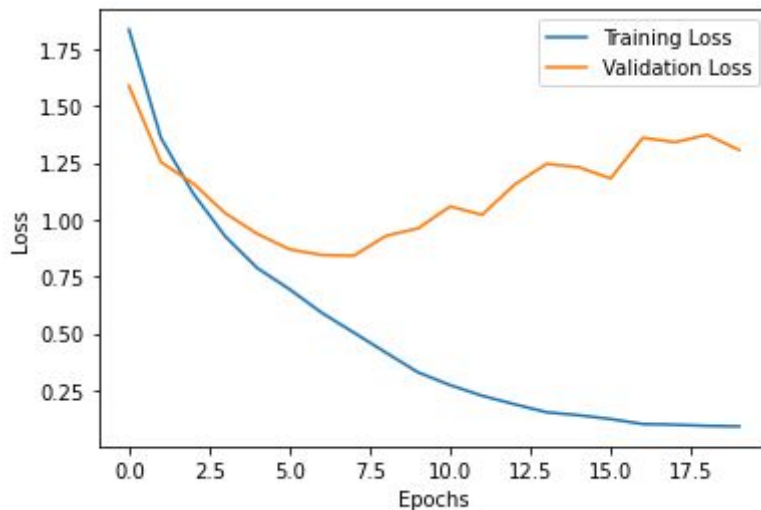

- The loss function chosen is Categorical Crossentropy, which is a standard loss function used for multi-class classification.
- Adam optimizer is used with a learning rate of 0.01.
- The learning rate scheduler used is CosineAnnealing.

```
criterion = nn.CrossEntropyLoss()
optimizer_rn = optim.Adam(model_rn.parameters(), lr=0.01)
exp_lr_scheduler = optim.lr_scheduler.CosineAnnealingLR(optimizer_rn, T_max=200)
```

- Resnet-18 model was used and trained from scratch for 20 epochs with batch size of 128.

```
model_rn = models.resnet18(pretrained=False)
num_fts = model_rn.fc.in_features
model_rn.fc = nn.Linear(num_fts, 10)
```

- Plots below show the Accuracies and Loss over epochs.



- Final performance of the model is given below:

```
Training Accuracy: 0.978
Validation Accuracy: 0.737
Testing Accuracy: 0.730
```

- The model took 6m 4s to train on Google Colab GPU.

Part 2

Goal

Initialize the ResNet-18 network with pretrained weights from ImageNet and then try to use these weights to improve the training for the CIFAR-10 dataset.

Transfer Learning:

We explore 2 different kinds of weight initialization techniques:

- Finetuning the CNN: Instead of random initialization, we initialize the network with a pretrained network (ImageNet). Rest of the training looks as usual.
- CNN as fixed feature extractor: Here, we will freeze the weights for all of the network except that of the final fully connected layer. This last fully connected layer is replaced with a new one with random weights and only this layer is trained.

Finetuning the CNN:



- Loaded the CIFAR-10 dataset, created a validation split, and applied normalization on the same lines as the previous problem.
- The loss function chosen is Categorical Crossentropy, which is a standard loss function used for multi-class classification.
- Adam optimizer is used with a learning rate of 0.001.
- Learning rate scheduler used is CosineAnnealing.

```
criterion = nn.CrossEntropyLoss()
optimizer_ft = optim.Adam(model_ft.parameters(), lr=0.001)
exp_lr_scheduler = optim.lr_scheduler.CosineAnnealingLR(optimizer_ft, T_max=200)
```

- Resnet-18 model was used which was pretrained on ImageNet dataset. The model was further trained for 20 epochs with batch size of 128.

```
model_ft = models.resnet18(pretrained=True)
num_fts = model_ft.fc.in_features
model_ft.fc = nn.Linear(num_fts, 10)
```

- Plots below show the Accuracies and Loss over epochs.



- Final performance of the model is given below:

```
Training Accuracy: 0.992
Validation Accuracy: 0.808
Testing Accuracy: 0.805
```

- The model took 6m 2s to train on Google Colab GPU.

CNN as fixed feature extractor:

- Loaded the CIFAR-10 dataset, created a validation split, and applied normalization on the same lines as the previous problem.
- The loss function chosen is Categorical Crossentropy, which is a standard loss function used for multi-class classification.
- Adam optimizer is used with a learning rate of 0.001.
- Learning rate scheduler used is CosineAnnealing.

```
criterion = nn.CrossEntropyLoss()
optimizer_ft = optim.Adam(model_ft.parameters(), lr=0.001)
exp_lr_scheduler = optim.lr_scheduler.CosineAnnealingLR(optimizer_ft, T_max=200)
```

- Resnet-18 model was used which was pretrained on the ImageNet dataset.
- The layers were freezed in 2 different configurations and the experiments are shown as follows:

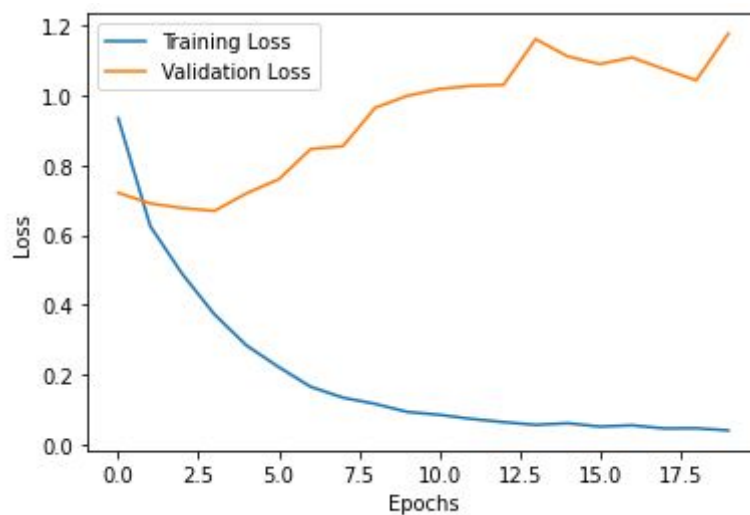

1. Model A:

- The first 6 Conv layers were freezed and the last fully connected layer is replaced with a new one with random weights and only this layer and the unfreezed layers are trained.

```
model_conv = torchvision.models.resnet18(pretrained=True)
num_fts = model_conv.fc.in_features
model_conv.fc = nn.Sequential(
    nn.Linear(num_fts, 256), nn.ReLU(), nn.Dropout(0.2),
    nn.Linear(256, 10))
```

```
count = 0
for name, child in model_conv.named_children():
    count += 1
    if count < 6:
        for name2, params in child.named_parameters():
            params.requires_grad = False
```

- Plots below show the Accuracies and Loss over epochs.



- Final performance of the model is given below:

```
Training Accuracy: 0.995
Validation Accuracy: 0.793
Testing Accuracy: 0.791
```

- The model took 5m 21s to train on Google Colab GPU.


2. Model B:


- The first 10 Conv layers were freezed and the last fully connected layer is replaced with a new sequential layer (as shown below) with random weights and only this layer and the unfreezed layers are trained.

```
model_conv = torchvision.models.resnet18(pretrained=True)
num_fts = model_conv.fc.in_features
model_conv.fc = nn.Sequential(
    nn.Linear(num_fts, 256), nn.ReLU(), nn.Dropout(0.2),
    nn.Linear(256, 10))
```

```
count = 0
for name, child in model_conv.named_children():
    count += 1
    if count < 10:
        for name2, params in child.named_parameters():
            params.requires_grad = False
```

- Plots below show the Accuracies and Loss over epochs.





- Final performance of the model is given below:

```

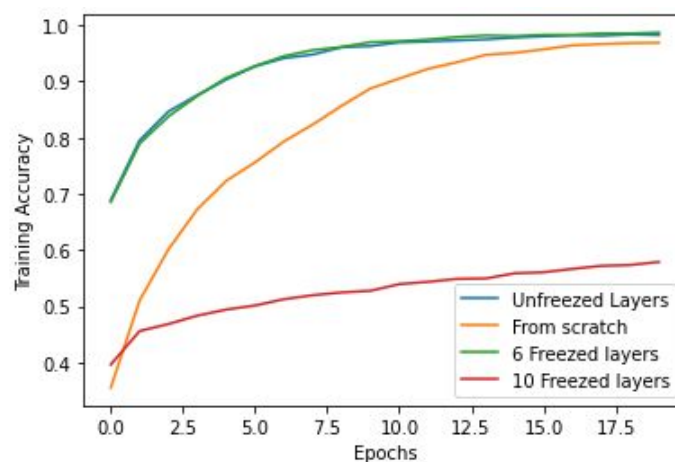
Training Accuracy: 0.615
Validation Accuracy: 0.492
Testing Accuracy: 0.491

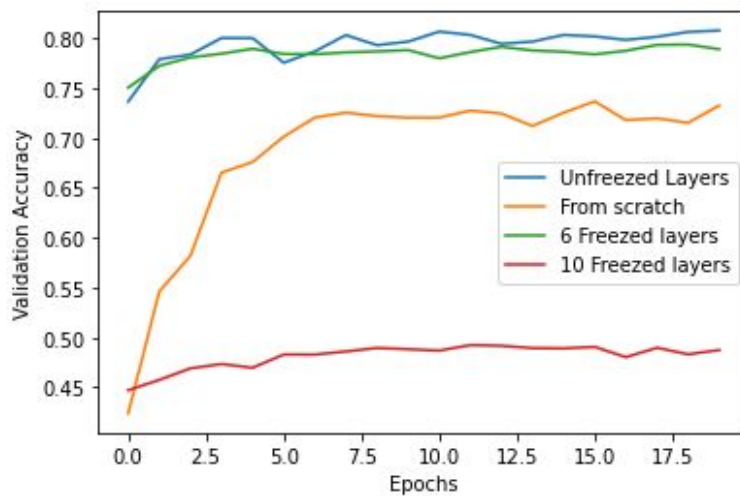
```

- The model took 3m 45s to train on Google Colab GPU.

Comparison

- The following plot shows the comparison between different Transfer learning configurations and base Resnet18 model.





- The following table compares the training time taken and final Test accuracy of above models for 20 epochs with a batch size of 128 .

<u>Model</u>	<u>Test Accuracy</u>	<u>Time Taken</u>
Resnet18 from scratch	0.73	6m 4s
Resnet18 pretrained (unfreezed)	0.805	6m 2s
Resnet18 pretrained (6 freezed layers)	0.791	5m 21s
Resnet18 pretrained (10 freezed layers)	0.491	3m 45s

- The Resnet18 pretrained (unfreezed) model gives the best Testing Accuracy but it takes more time.
- The Resnet18 pretrained (10 freezed layers) takes the least time but gives significantly poor performance.
- Therefore in terms of both Test accuracy and time taken to train, Resnet18 pretrained (10 freezed layers) performs the best.

Part 3

Goal

Train the network from scratch with Tiny-CIFAR-10 dataset using data augmentation techniques to improve the performance.

Procedure

- Loaded the CIFAR-10 dataset, created a split to create mini CIFAR-10 dataset.
- Data Augmentation techniques used:
 - Resize
 - Random Horizontal Flip
 - Random Rotation
- The following code snippet shows the mini CIFAR-10 dataset creation and data augmentation transformation.

```
train_transform = transforms.Compose([
    transforms.Resize(224),
    transforms.RandomHorizontalFlip(p=.40),
    transforms.RandomRotation(30),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])])

test_transform = transforms.Compose([
    transforms.Resize(224),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])])

dataset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=train_transform)
train_ds, val_ds = random_split(dataset, [5000, 45000])
test_ds = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=test_transform)
```

- The loss function chosen is Categorical Crossentropy, which is a standard loss function used for multi-class classification.
- Adam optimizer is used with a learning rate of 0.001.
- Learning rate scheduler used is CosineAnnealing.

```
criterion = nn.CrossEntropyLoss()
optimizer_ft = optim.Adam(model_ft.parameters(), lr=0.001)
exp_lr_scheduler = optim.lr_scheduler.CosineAnnealingLR(optimizer_ft, T_max=200)
```

- Resnet-18 model was used and trained from scratch.


Introducing Dropout:

- To introduce regularization a Dropout layer was added as shown below:

```
model_ft = models.resnet18(pretrained=False)
num_fts = model_ft.fc.in_features

model_ft.fc = nn.Sequential(
    nn.Dropout(dropate),
    nn.Linear(num_fts, 10))
```



- The dropout rate was varied from 0.1 to 0.4 and Training and Testing Accuracy for the respective models is given by the graph below.



- Dropout rate of 0.3 performed the best as shown by the graph.

Final Model:

- The final model with Dropout rate of 0.3 was trained for 50 epochs with batch size of 128.
- Plots below show the Accuracies and Loss over epochs.




- Final performance of the model is given below.

```
Training Accuracy: 0.896
Validation Accuracy: 0.728
Testing Accuracy: 0.728
```


- The model took 23m 51s to train on Google Colab GPU.

Part 4


- We visualise various activation layers of ResNet-18 at different levels.
- Activations are basically the pixel values that have high correlation with correctly identifying what an object is.
- This basically gives us a way to visualize what each layer is doing and what it's trying to focus on within a certain image that is passed for classification.
- The figure of a ship shown below to make the analysis.



- These are two of the initial level filters (*conv1_conv* and *conv2_block1_1_conv*) mapped for the ship given as an input. As can be seen, it gives a rough idea of the shape of the ship and the contents on the image around it like the sky and sea.




- The activation maps change significantly for later layers. The following plot corresponds to the 25th layer (*conv2_block2_3_conv*) of the model.
- As can be seen, the feature maps have considerably become less complicated than their initial counterparts and it is difficult quantitatively to understand what they try to depict.




- Since CIFAR-10 has its images in a 32x32 pixel quality, we don't get the most accurate activation visualization.
- This experiment can be run again with images of a higher quality (more pixels) to see a better representation of how the activations look for a particular image.

- Another example containing an image of a cat has been discussed below.



- First 2 layers have been shown below.



- 25th layer.

